

## Presentation 04

Objective: generate and present credit values predictions using RFR

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_absolute_

In [2]: df = pd.read_csv('/home/gmelao/Desktop/default-of-credit-card-clients.csv')
df.columns = df.iloc[0]
df.drop(0, inplace = True)
df.set_index('ID', inplace = True)
pd.set_option('display.max_columns', 24)
pd.set_option('display.max_rows', 24)

In [3]: df = df.apply(lambda df: pd.Series(map(float, df)))
```

### Train Test Split

Split arrays or matrices into random train and test subsets.

```
In [4]: X = df.drop('LIMIT_BAL', axis=1)
y = df['LIMIT_BAL']

In [5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_st
```

### Random Forest Regressor

A random forest is a meta estimator that fits a number of classifying decision trees on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting. The sub-sample size is controlled with the max\_samples parameter if bootstrap=True (default), otherwise the whole dataset is used to build each tree.

```
In [6]: model = RandomForestRegressor()

In [7]: model.fit(X_train, y_train)

Out[7]: ▼ RandomForestRegressor
RandomForestRegressor()

In [8]: preds = model.predict(X_test)
preds

Out[8]: array([ 71300., 180600., 131500., ..., 235100., 168900., 116800.])

In [9]: pd.DataFrame(preds, columns=['Credit Prediction'])
```

Out[9]:

Credit Prediction	
0	71300.000000
1	180600.000000
2	131500.000000
3	154300.000000
4	140100.000000
...	...
9895	154716.666667
9896	140900.000000
9897	235100.000000
9898	168900.000000
9899	116800.000000

9900 rows × 1 columns

## Metrics

MAE: it is the mean of the absolute error. This gives less weight to outliers, which is not sensitive to outliers.

MSE: MSE is a combination measurement of bias and variance of your prediction

RMSE: Take a root of MSE would bring the unit back to actual unit, easy to interpret your model accuracy.

MAPE: MAPE is the sum of the individual absolute errors divided by the demand (each period separately). It is the average of the percentage errors.

```
In [11]: mae = mean_absolute_error(y_test, preds)
```

```
In [12]: mse = mean_squared_error(y_test, preds)
```

```
In [13]: rmse = np.sqrt(mean_squared_error(y_test, preds))
```

```
In [14]: mape = mean_absolute_percentage_error(y_test, preds)
```

```
In [15]: print(f'MAE: {mae}')
print(f'MSE: {mse}')
print(f'RMSE: {rmse}')
print(f'MAPE: {mape}')
```

```
MAE: 68484.06770094384
MSE: 9025328793.808342
RMSE: 95001.7304779673
MAPE: 0.7853820680953061
```

```
In [ ]:
```

```
In [17]: y_train.mean()
```

Out[17]: 167197.61194029852

```
In [20]: baseline = np.arange(9900)
baseline.fill(y_train.mean())
```

```
In [21]: mae_baseline = mean_absolute_error(y_test, baseline)
mse_baseline = mean_squared_error(y_test, baseline)
rmse_baseline = np.sqrt(mean_squared_error(y_test, baseline))
mape_baseline = mean_absolute_percentage_error(y_test, baseline)
```

```
In [22]: print(f'MAE: {mae_baseline}')
print(f'MSE: {mse_baseline}')
print(f'RMSE: {rmse_baseline}')
print(f'MAPE: {mape_baseline}')
```

MAE: 105694.56606060606  
MSE: 17081185206.624243  
RMSE: 130695.00834624191  
MAPE: 1.669181588976168

In [ ]:

```
In [25]: print(f'MAE / MAE_BASELINE: {mae_baseline/mae}')
print(f'MSE / MSE_BASELINE: {mse_baseline/mse}')
print(f'RMSE / RMSE_BASELINE: {rmse_baseline/rmse}')
print(f'MAPE / MAPE_BASELINE: {mape_baseline/mape}')
```

MAE / MAE\_BASELINE: 1.5433453299262683  
MSE / MSE\_BASELINE: 1.8925831509143989  
RMSE / RMSE\_BASELINE: 1.3757118706016893  
MAPE / MAPE\_BASELINE: 2.1253115608104425

In [ ]: